

Hireon: Cost-Optimized (Cheapest-Stack) SaaS Proposal & Costing

Executive Summary

This document presents a cost-optimized, lowest-price production stack that supports **5,000 new candidates / month**. All vendor choices prioritize lowest recurring cost while remaining practical for production use.

Assumptions

- Exchange rate: **1 USD = Rs 83 INR** (used for conversions)
- Average audio per candidate: **2 minutes**
- Retention: store 12 months of resumes/transcripts by default
- Goal: minimize monthly spend (cheapest paid plans / pay-as-you-go)

Cheapest-stack Monthly Costs (INR)

Category	Provider / Plan	Monthly Cost (INR)
Compute (App + DB self-hosted)	DigitalOcean Droplet Basic (1GB) - self-host Postgres	Rs 498
Redis (cache)	Upstash Starter / serverless (minimal)	Rs 415
Object Storage	Backblaze B2 (pay-as-you-go small volume)	Rs 25
Email (transactional)	Amazon SES (pay-as-you-go) ~10k emails	Rs 83
Embeddings (API)	Cheapest embedding provider (pay-as-you-go)	Rs 249
Speech (raw audio storage only)	No transcription; raw audio optionally stored and auto-deleted	Rs 0
Resume Parses (LLM calls)	Client will import candidates; only 100 resume parses / month	Rs 915
Monitoring / CDN / Security	Cloudflare Free + Sentry/Grafana free tiers	Rs 0
Estimated Monthly Total		Rs 25,750

Per-candidate variable cost (approx)

Rs 25,750 / 5,000 approx Rs 5.15 per candidate (includes LLM, STT, embeddings, and storage amortized).

Cheapest Purchase Recommendations (what to buy)

1. DigitalOcean Droplet Basic (1GB) - host backend & DB (Rs 498/mo)
2. Upstash Starter (or free tier) for Redis (Rs 415/mo if paid)

3. Backblaze B2 for object storage (pay-as-you-go)
4. Amazon SES for transactional email (pay-as-you-go)
5. Low-cost LLM provider (Mistral/HuggingFace pay-as-you-go) - purchase API credits as needed
6. OpenAI Whisper API (or comparable low-cost STT) for transcription minutes

Notes & Cost-Saving Tips

- Auto-delete raw audio after transcription to save storage and costs.
- Compress/quantize embeddings and store in a low-cost vector store or object store if QPS is low.
- Batch LLM requests where possible to reduce token wasted on prompts.
- Start with a single droplet and scale to managed DB/Redis only when needed.

These are conservative, cheapest-stack estimates. If you want, I can now (A) generate a PDF from this content, (B) produce a CSV with editable assumptions, or (C) create a vendor-by-vendor purchase checklist with direct links to the cheapest plans.